


Top 10 películas (rating por promedio de reseñas)



IMDB - Model Results


Total reviews in DB: 5463

Reviews used for modeling (first 4000): 3997

Reviews used for PDF visuals/evaluation (JSON): 2013


Models trained: Naive Bayes, Linear SVM, Decision Tree

Top TF-IDF terms (reviews from JSON)




Distribución de polaridad (reseñas JSON)


Número de reseñas



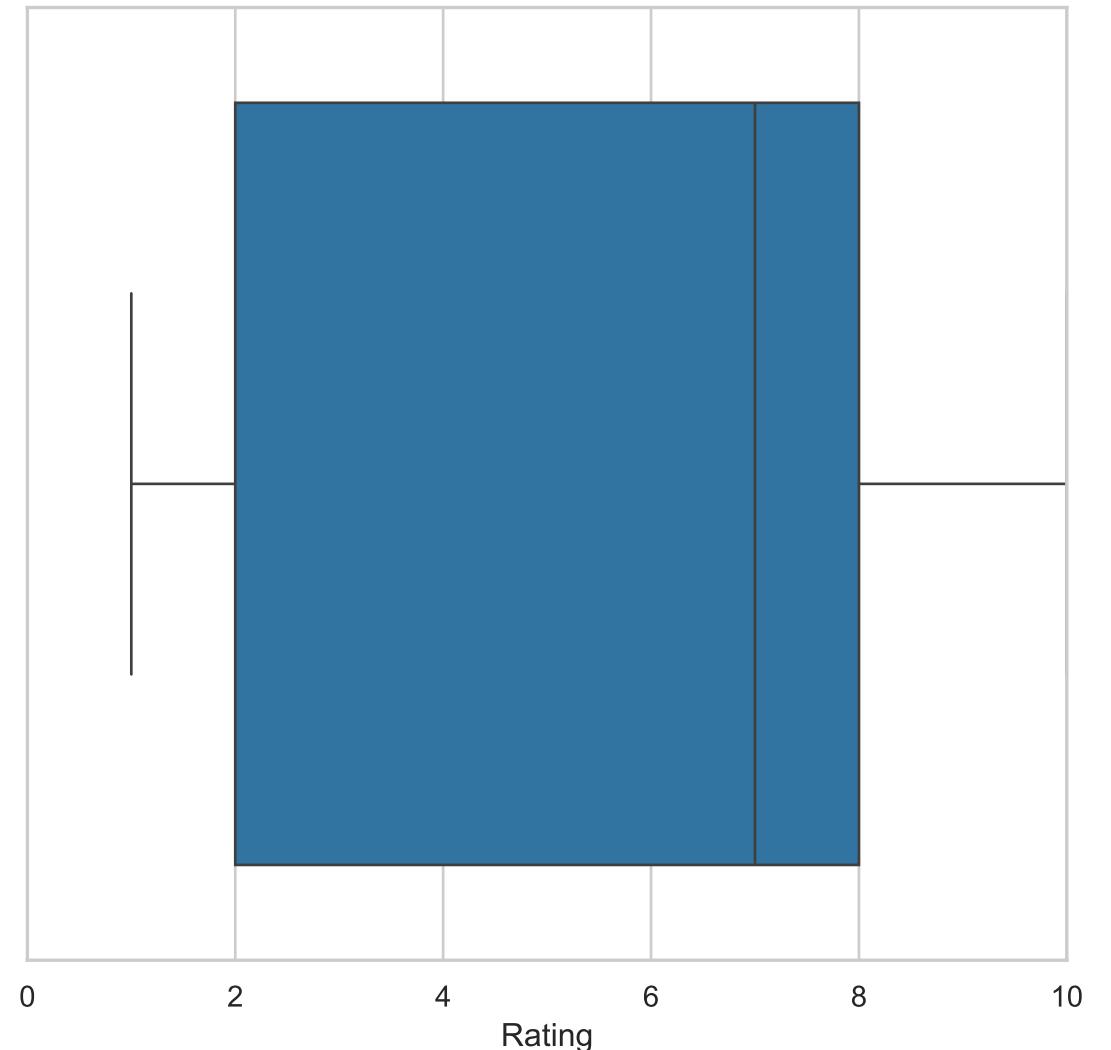
Porcentaje de polaridad (reseñas JSON)




Polaridad por película (top 10 por #reseñas en JSON)



Distribución de calificaciones (boxplot) - global



Distribución de calificaciones por polaridad (boxplot)



Distribución de calificaciones por película (top 8 por #reseñas)

Película

The Godfather

The Godfather: Part II

12 Angry Men


Trois couleurs: Blanc

Goodfellas


On the Waterfront

The 39 Steps


The Dark Knight




KMeans clusters (TruncatedSVD 2D) - JSON reviews




Cluster 0: 480 pts
 Cluster 1: 395 pts
 Cluster 2: 444 pts
 Cluster 3: 319 pts
 Cluster 4: 92 pts
 Cluster 5: 192 pts
 Cluster 6: 91 pts




Cluster 0 - top TF-IDF words (centroid)




Cluster 1 - top TF-IDF words (centroid)




Cluster 2 - top TF-IDF words (centroid)




Cluster 3 - top TF-IDF words (centroid)




Cluster 4 - top TF-IDF words (centroid)




Cluster 5 - top TF-IDF words (centroid)




Cluster 6 - top TF-IDF words (centroid)




Naive Bayes - Confusion Matrix




Linear SVM - Confusion Matrix




Decision Tree - Confusion Matrix




ROC Curves (JSON eval)



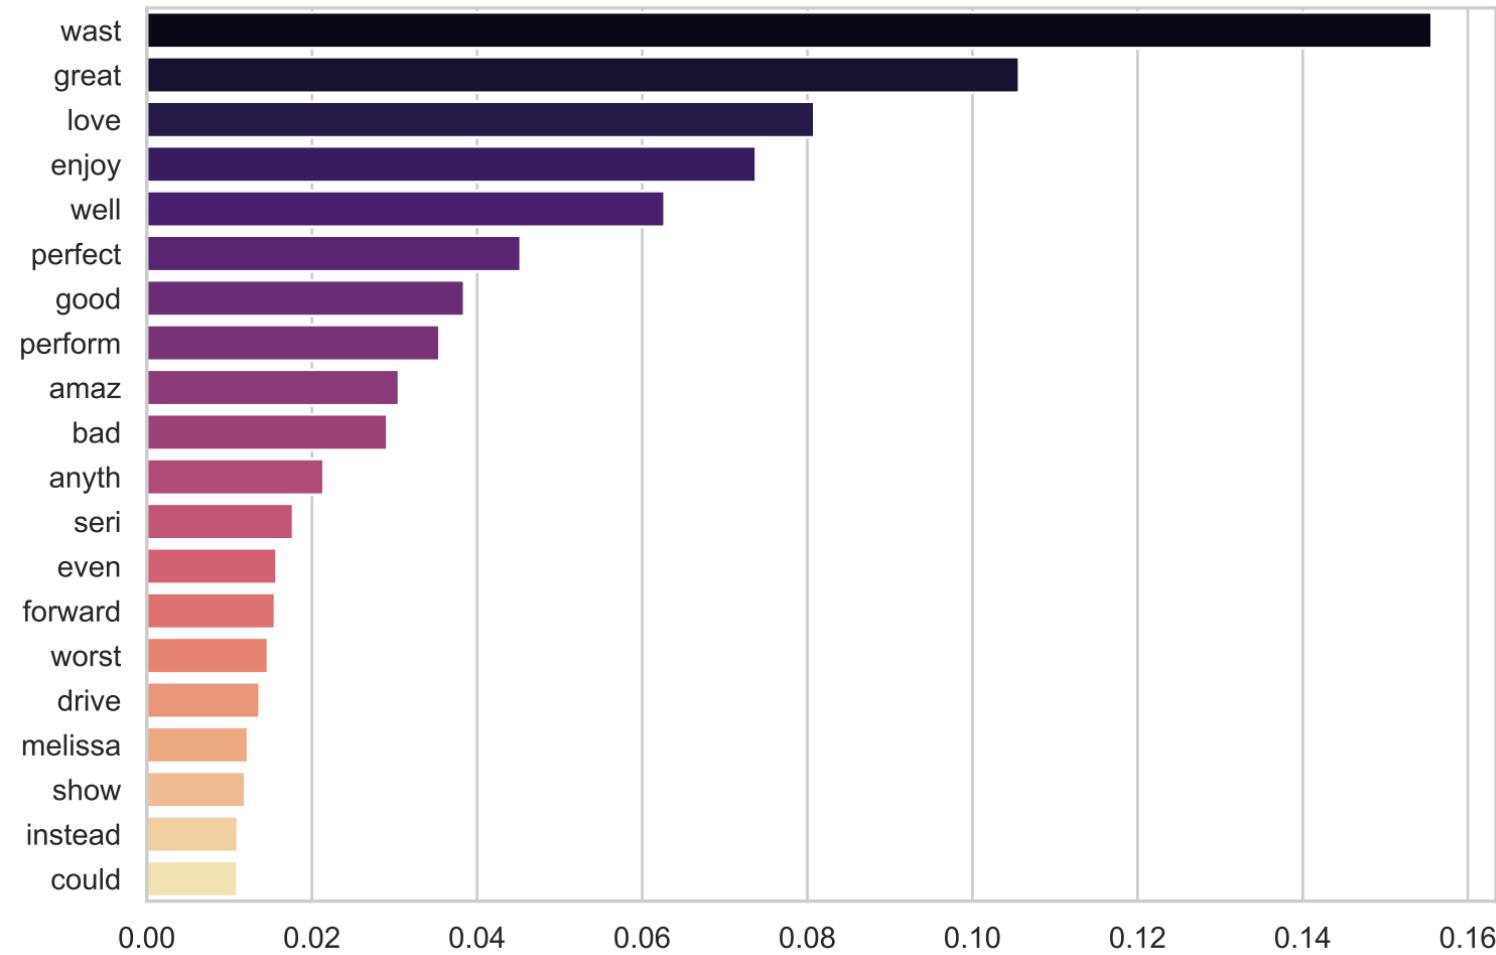
Precision-Recall Curves (JSON eval)




Model metrics (evaluated on JSON reviews if available)




Decision Tree - top feature importances (training)




LDA Topic 0 - top words (LDA)




LDA Topic 1 - top words (LDA)




LDA Topic 2 - top words (LDA)



LDA Topic 3 - top words (LDA)



LDA Topic 4 - top words (LDA)



Classification reports (evaluated on JSON reviews if available)

Model: Naive Bayes

	precision	recall	f1-score	support
NEG	0.93	0.73	0.82	1005
POS	0.78	0.94	0.85	1008
accuracy			0.84	2013
macro avg	0.85	0.84	0.84	2013
weighted avg	0.85	0.84	0.84	2013

Model: Linear SVM

	precision	recall	f1-score	support
NEG	0.92	0.84	0.88	1005
POS	0.85	0.93	0.89	1008
accuracy			0.88	2013
macro avg	0.89	0.88	0.88	2013
weighted avg	0.89	0.88	0.88	2013

Model: Decision Tree

	precision	recall	f1-score	support
NEG	0.72	0.80	0.76	1005
POS	0.77	0.69	0.73	1008
accuracy			0.75	2013
macro avg	0.75	0.75	0.74	2013
weighted avg	0.75	0.75	0.74	2013